

How Does Content Drive Viewership?

Dave Holtz¹, Jeremy Yang², Michael Zhao^{3*}

Abstract

Why do some webpages receive massive numbers of page views? To determine how content drives viewership, we construct a unique dataset of all articles published by the New York Times (NYT) in August 2013. Our dataset is built from 2 major components, the NYT’s internal web traffic data and web content data extracted from the NYT website. We use the internal web traffic data to accurately track the number of page views of each article as well as construct a set of robust control variables such as the desk and section of each article. To build textual content features, we use various machine learning and statistical natural language processing techniques on our extracted content data to produce perplexity scores and to determine the sentiment and reading difficulty of each article. We also generate indicators that denote the presence of pictures, videos, etc. Lastly, we use some secondary data sources to build some additional control features such as estimated author gender and author popularity. We take all these features and build a predictive regression model. Overall, we do find that our textual features improve our predictive power, though rather modestly. However, due to time constraints, there remain many textual features we haven’t yet implemented. We believe that adding these additional features may offer more significant improvements in performance.

¹dholtz@mit.edu

²zheny@mit.edu

³mfzhao@mit.edu

**Michael is using this project to fulfill Final Project requirements for both 6.864 as well as 6.867. Dave and Jeremy are only using this project for 6.867.*

Contents		1. Introduction	
1	Introduction	1	In today’s digital economy, both individuals and companies are very interested in attracting users to visit their websites in order to earn ad revenue. While many factors might motivate a user to visit a particular page, certainly one important factor is the content contained within that webpage. This paper explores the relationship between the content of a webpage and the number of page views it receives by constructing a unique dataset of all articles published by the New York Times (NYT) during August 2013. This dataset is built from two major components: the NYT’s internal web traffic data and webpage content data parsed from the NYT website.
2	Data	2	
2.1	NYT Internal Web Traffic Data	2	
2.2	Parsed NYT Webpage Content Data	3	
3	Constructed Features	4	
3.1	Article Sentiment	5	
3.2	Article Perplexity	7	
4	Predictive Regression Model	9	
4.1	The Effect of Textual Features	11	
5	Future Work	13	
6	Division of Labor	14	
	Acknowledgments	14	Typically, a study such as ours tends to be very difficult to conduct accurate measures of viewer-
	References	14	

ship are either private or unavailable¹. Even in cases where page views are publically, for example Youtube or various other video sharing sites, feature extraction of the content is far too challenging given the tools we have available to us today². Fortunately, our access to the the NYT's internal web traffic data allows us to exactly measure the number of page views an article receives. The web traffic data is rather rich and also includes internal meta-data that we use to build various control features. Moreover, since we are working with mostly textual data, we are able to take advantage of recent advancements in machine learning and statistical NLP to do feature extraction on article text.

A similar study by Berger and Milkman (2012) [1] examines the relationship between content and word-of-mouth virality. They find that the emotional content of a NYT article is predictive of its virality. Using simple measures of an article's sentiment and emotionality, Berger and Milkman show that positive articles are more likely to show up on the New York Times "Most-Emailed" list. They also show that articles that evoke high physiological positive or negative arousal (such as awe or anger) tend to be more viral than articles that evoke deactivating emotions (sadness). We build on this study in two ways: first, we relate an article's content back to the number of page views it receives rather than its virality³. Second, we employ more sophisticated machine learning feature extraction techniques.

2. Data

2.1 NYT Internal Web Traffic Data

Our NYT internal web traffic dataset is a record of all individual user activity on the NYT website covering the period of April 3rd, 2013 to October

31st, 2013. Each time a user⁴ moves from one page to another on the NYT website, this activity is captured as an individual line json. This data is incredibly detailed and can potentially tell us who accessed what page when and from what location. In addition it tracks the duration a user stays on a webpage as well as the path a user took to get to her current webpage.⁵ Overall, this data is over 20 terabytes in size and contains over 3 billion page views. It's important to note that not all page views are page views of content we care about. For example, some events that are also tracked in the data are searches or changes in user account settings. Since the scope of this dataset is so large, we initially restrict this project to a single month, August 2013, with plans of extending our analysis to our entire dataset for future work. We further restrict our dataset to only include articles or blogposts since these are the pieces of content that have contain readily extractable text data. For the sake of brevity, we use the term "articles" to refer to both articles and blogposts unless we explicitly state otherwise.

We make a basic first pass through the data to simply to obtain a list of urls. After cleaning up the url data to ensure each url mapped uniquely to a particular piece of content, we were left with a total of 6682 urls. We then parse all the web data for the month of August and the first week of September count the number of impressions a url receives. In order to make an apples to apples comparison between articles we only count the number of page views received up to 7 days after publication since an article that's been out longer should have more page views in expectation. Given the tendency for the viewership of an article to drop off sharply soon after publication (as recency is important factor in determining of news readership), our 7-day mea-

¹While precise viewership data tends to be not available openly, oftentimes researchers use related observables, such as Facebook likes, or number of Tweets and Retweets

²Though this is quickly changing with advances in computing power and development of better and better machine learning methods

³Which companies arguably care more about since word-of-mouth virality is usually a means to increase page views

⁴in this case, a "user" is uniquely defined by a device/browser id. So while the same person might have multiple devices or may use multiple browsers, the NYT backend treats each device/browser combination as a unique "user" even though in reality its all the same person. In some cases, we are able to link various id's together if the person happens to register an official user account on the NYT website and then logs into her account from multiple devices/browsers.

⁵Big brother is watching you and we are big brother. Welcome to the age of Big Data

sure generally represents the vast majority (well over 90%) of total page views that an article receives⁶. Even after all this subsampling, our data still consists of 248,161,455 page views⁷. As seen in Figure 1 below, the distribution of page views is highly skewed with very heavy tails. After applying a log transformation (as seen in Figure 2), our distribution looks considerably more normal.




Figure 1. Histogram of Articles by Number of Page Views

Table 1. Page Views Distribution Summary Statistics

Min	1
Max	2545288
Mean	37138.8
Median	10298.5
StDev.	88972.9
Skew	9.52191
Kurt.	173.061
Obs	6682

In addition to aggregating the counts, when we parse the internal web traffic data, we make sure to extract various relevant article meta-data features such as the headline, time of publishing, authors, section, desk, and the NYT's internal article content description tags (if available).

⁶at least for a reasonable stretch of time

⁷Though one video by PSY completely crushes this number




Figure 2. Histogram of Articles by Log of Page Views

Table 2. Log Page Views Distribution Summary Statistics

Min	0
Max	14.74975
Mean	9.122868
Median	9.239754
StDev.	2.028668
Skew	-1.270368
Kurt.	3.800911
Obs	6682

2.2 Parsed NYT Webpage Content Data

Unfortunately, the NYT internal web traffic data does not contain the actual content displayed on each webpage, which is a very important aspect of our project. Luckily, all this content is freely hosted on the NYT website! Although the NYT limits the number of free articles you can access per month, the tracking system is cookie based. This means that scraping the raw html content via wget or something similar does not contribute to your article count limit⁸. Specifically, we used the python library “newspaper” to download the html content from the NYT website and then extract all the raw textual data from the html. We applied some additional regular expressions filters to clean up what the library missed. In total, our 6682 articles contain 4,685,021 words of text. We find that the distribu-

⁸Alternatively, you can just keep clearing your cookies

tion of article length, much like the distribution of page views, is highly skewed and heavy tailed. We again apply a log transformation. Histograms and various summary statistics of for the distribution of article length and log article length are provided below:




Figure 3. Histogram of Articles by Word Count

Table 3. Log Page Views Distribution Summary Statistics

Min	7
Max	8941
Mean	701.1405
Median	625
StDev.	591.9535
Skew	3.335127
Kurt.	24.35949
Obs	6682

In addition to extracting the raw text data, we made sure to check if the articles also contained any additional non-textual content such as pictures or videos. We created indicator variables that denote the presence of such content within an article.

3. Constructed Features

Using our collected this data, as well as some data from some additional secondary sources we construct the features that make up our predictive regression model. These features include the Flesch reading ease, the estimated gender of the author(s),




Figure 4. Histogram of Articles by Log Word Count

Table 4. Log Page Views Distribution Summary Statistics

Min	1.945910
Max	9.098403
Mean	6.181044
Median	6.437751
StDev.	1.004414
Skew	-1.178993
Kurt.	1.944611
Obs	6682

the popularity of the author(s), variables indicating the section the article appeared in and the article's content type, the sentiment of the article text, and the perplexity of the article text. We provide a full list of these features below, as well as the methodology used to construct them. Where appropriate, we include discussion of testing and validation of our features and our algorithms.

Flesch Reading Ease One can conceive of a few competing hypotheses that relate the readership of content to the ease with which people can read it. Maybe more complicated pieces of text are more engaging, and are more likely to be read. On the other hand, perhaps pieces of text that are easier to read will be consumed by more people. In order to capture relationships such as these in our data, we include the Flesch reading ease. The

Flesch reading ease is a metric developed by Flesch in 1948 [2]. The score indicates how difficult a piece of English text is to understand. Lower scores correspond to more difficult passages, and the highest score attainable is 120.0. The formula for calculating a passage's Flesch reading ease is

$$206.835 - 1.015 \left(\frac{\# \text{ words}}{\# \text{ sentences}} \right) - 84.6 \left(\frac{\# \text{ syllables}}{\# \text{ words}} \right) \quad (1)$$

To calculate the Flesch reading ease, we use the python library “textstat.” Despite the fact that the above formula is relatively straightforward, the task of counting the number of syllables in a block of text is non-trivial, so we rely on “textstat” to do so accurately. In cases where the Flesch reading ease was for some reason null (e.g., a blog post containing only a picture), we assign the Flesch reading ease its median value.

Author Popularity We attempt to include some measure of a particular author's popularity. It stands to reason that a new article by Paul Krugman or A.O. Scott should garner more readership than a new article by an unknown graduate student enrolled in 6.867 at MIT!

In order to measure something that will serve as a decent proxy for popularity, we programmatically searched for every distinct author in our dataset on Bing and recorded the number of search results that were returned by the query. In cases where a particular article has more than one distinct author, we calculate an “effective” popularity by simply averaging number of search results over all article authors.

Author Gender While we certainly don't think that an author's gender has a causal impact on the readership on an article. We believe that this feature allows us to control for some

latent unobserved heterogeneity. For example, it's not a stretch to think that the experiences of women authors are drastically different from male authors and these differences are reflected in their writing. We construct this feature that indicates the most likely gender of the article author(s). In cases where the gender of the author is unclear (e.g., Robin) or there are likely multiple authors with different genders (e.g., The New York Times Staff), we record a third gender value, “ambiguous / unknown.”

Our gender data is gathered by cross-referencing the first names of all of the authors in our dataset against U.S. Social Security Administration baby name data from 1935 to 1997. If over 90% of the babies with a given name have been male, we assume a given author is male. If over 90% of the babies with a given name have been female, we assume a given author is female. Otherwise, we record “ambiguous / unknown.”

Material Type, Section, Desk, and Article Type

We also include dummy variables including the material type (e.g., ‘News’ or ‘Obituary’), publishing desk (e.g., ‘Weekend’ or ‘Real Estate’), article type (‘Blog post’ or ‘Article’), and section (e.g., ‘Movies’ or ‘World’). The hypothesis driving the decision to include these variables is that certain types of content (e.g., political news or international affairs) may be more widely read than local material (such as real estate) or less popular sections of the NY Times (e.g., the sports section).

We also build features that attempt to capture the article sentiment and the article text perplexity. Since the construction of these features was considerably more complex and involved validation of our algorithms, we discuss these two features in separate subsections.

3.1 Article Sentiment

In order to measure article sentiment, we use a Naïve Bayes text classification algorithm, as described in Rennie et al (2003) [3]. We assume that

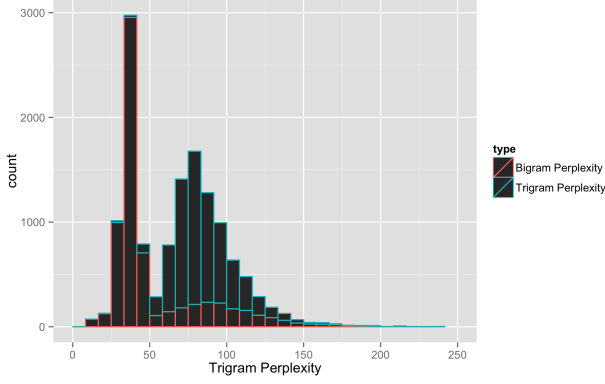


Figure 5. Histogram of Bing Search Results




Figure 6. Histogram of log(Bing Search Results)

each article in our corpus can belong to one of three classes - ‘negative’ sentiment, ‘neutral’ sentiment, and ‘positive’ sentiment, which we will denote as C_k . The Naive Bayes model assumes that the likelihood of observing a given article $\mathbf{x} = (x_1, \dots, x_n)$, where x_i is the number of times that word i appears, is

$$p(\mathbf{x}|C_k) = \frac{(\sum_i x_i)!}{\prod_i x_i!} \prod_i p_{ki}^{x_i}, \quad (2)$$

where p_{ki} is the probability a word w_i conditional on a document of being of class K . Applying a log transformation, we can compute $\log(p(\mathbf{x}|C_k))$ as:

$$\log(p(\mathbf{x}|C_k)) = \log(p(C_k)) + \sum_{i=1}^n x_i \cdot \log(p_{ki}). \quad (3)$$

We coded up a basic implementation of the Naive Bayes algorithm, drawing heavy inspiration

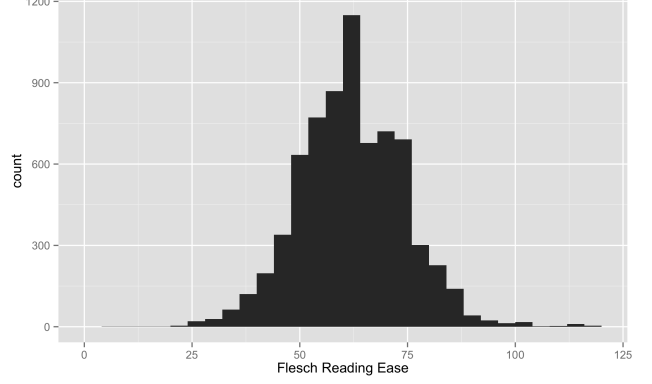


Figure 7. Histogram of Flesch Reading Ease

from Greg Lamp’s 2014 python tutorial on Naive Bayes [4]. In order to get the probabilities $p(C_k)$ and p_{ki} , we needed some labeled training data. In order to obtain these labels, we selected a random subset of 200 articles from our dataset and created a task on Amazon Mechanical Turk. Each Turker was asked to score the sentiment toward the subject of the article in question on a scale from -2 to +2, with -2 being extremely negative and +2 being extremely positive. In order to make sure these scores were relatively robust, we recorded 5 scores for every article from 5 different Turkers and calculated the average sentiment score. We classified any article having an average score greater than 0.5 as ‘positive’. Any article with an average sentiment less than -0.5 was ‘negative.’ Any other articles were classified as ‘neutral.’ Ultimately, our labels were 66% neutral, 14.5% negative, and 19.5% positive. This is unsurprising, as a newspaper such as the New York Times likely strives for neutrality when reporting on most topics.

We wanted to determine how our Naive Bayes implementation did compared to an off-the-shelf implementation of the same algorithm. In order to do so, we trained NLTK’s multinomial Naive Bayes classifier on the same training data, and then evaluated the sentiment of 1,000 articles. A comparison of the two implementations is found below, where the columns indicate the prediction by the NLTK Naive Bayes implementation and the rows indicate the prediction by our implementation:

	negative	neutral	positive
negative	0	54	0
neutral	2	894	0
positive	0	50	0

Overall, we find 89.4% agreement between the two algorithms. However, alarmingly, the NLTK implementation seems to predict neutral an overwhelming percentage of the time (99.8% of the time). This warrants further investigation, and may be due to small differences in implementation, or peculiarities in the sample of 1,000 articles we chose to compare the two algorithms across. In any case, our algorithm seems to be performing in the same neighborhood as the NLTK implementation (if not better), so we feel relatively comfortable moving forward using our sentiment labels.

3.2 Article Perplexity

In order to determine the perplexity score, we first need to build some language model that gives us the probabilities of each word. While perplexity typically is a measure of how well a probability distribution can predict a sample, in our context, we interpret perplexity essentially as a measure of article “uniqueness”. The argument here is that if our language model can’t predict the language used in article very well, then the language used in the article is atypical relative to the corpus used to build the language model. Hence, given some language model, an article’s perplexity is given by:

$$2^{-n \cdot \sum_{i=1}^n \log p(w_i)} \quad (4)$$

where n is the length of the article, and $p(w_i)$ is the probability of the i -th word in the article. We think that perplexity might have some predictive power since people generally have a preference for novelty. If many news articles about the same story are all using highly similar language, an article that covers the story using atypical language is likely unique in some way or another which may drive people to read it more or less. For this section we generally follow the 6.864 Lecture Notes [5].

As for our paper, we construct a couple of different language models. First, we build a simple bi-gram language to use as a baseline. We also build more sophisticated word vector based n-gram neural network language models.

We split our articles into training (70%), validation (15%), and test corpora (15%). In order to keep the size of our vocabulary relatively manageable, we ignore any case sensitivities (so “Cat” is the same as “cat”). Furthermore, we only include a word if it appears at least 5 times. Words that don’t make this cutoff are mapped to a generic “rare word” indicator. Lastly, we also map any numbers (that is numbers comprised of digits, not numbers written with words) to a generic “number” indicator. Ultimately, this leaves us with a vocabulary size $|V|$ of 29359. To estimate a bigram model, we simply need to compute the counts in our training corpus. Specifically, the probability of some word w_i conditional on its preceding word w_{i-1} is given by:

$$p(w_i|w_{i-1}) = \frac{\text{count}(w_{i-1}, w_i)}{\text{count}} \quad (5)$$

However, it’s reasonable to expect that there might be bigrams in the development or test corpora that aren’t observed in the training corpus considering that for our vocabulary size, there are nearly 900 million unique bigrams and our training data contains just over 3 million observations. If this is the case, then any article with a bigram unobserved in the training corpus would be assigned a predicted likelihood of 0. Needless to say, this is very bad. In order to avoid this issue, we apply a technique called add- α smoothing. Add- α smoothing just adds α to each cell in the probability table. Hence, after smoothing, no bigram, given a fixed vocabulary V , will ever have 0 probability. This changes our estimated word probabilities to:

$$p(w_i|w_{i-1}) = \frac{\text{count}(w_{i-1}, w_i) + \alpha}{\text{count} + \alpha|V|} \quad (6)$$

Essentially, this smoothing is acting like a probabilistic Robin Hood, it takes probability mass from the “rich” and distributes it to the “poor”. Additionally,

this smoothing actually has a nice Bayesian interpretation as well. For the bigram language model case, we can say that words follow a multinomial distribution conditional on its preceding word. The smoothed language model is equivalent to the posterior of horizontally stacked multinomial likelihoods with a symmetric Dirichlet prior with a parameter value of α .

We use our validation set to determine the optimal value of α by seeing what value of α minimizes the negative average log-likelihood per word (NALL) of our validation corpus:




Figure 8. Negative Average Log-Likelihood per Word as a Function of α

Here we see exactly what we should expect. The NALL on the training corpus is lower than the NALL on the validation corpus. Furthermore, the NALL on our training corpus monotonically increases as we increase α while the NALL on the validation corpus initially decreases hits a minimum, and starts to increase. For our language model, the optimal α is 0.004. This level of smoothing produces the following NALLs for our 3 corpora:

Table 5. Negative Average Log-Likelihood for Smoothed Bigram Language Model

Training	Validation	Test
4.98487	6.42916	6.41888

One criticism of standard n-gram language models is that they are rather sensitive to the training data. Suppose that two words generally have fairly

interchangeable use cases (that is the sets of words that tend to precede these words tend to be highly similar). For example, consider “coffee” and “tea”. Further suppose that for whatever reason, the phrase “drink coffee” is predominantly featured in the training data but “drink tea” is not. Then a standard bigram language model would assign a very low probability to “tea” if it were preceded by the word “drink” even if the in the training data contains many instances in which they are fairly interchangeable (“buy coffee/tea”, “brew coffee/tea”, etc.).

Hence, we build word vector n-gram neural network models to see if we can achieve better performance. By using dense word vectors rather than one-hot encodings to represent words, we’re able capture the underlying similarity of words and their meanings. For example, “coffee” is more similar to “tea” than it is to “car”. Rather astonishingly, these word vector values can be trained through back propagation just like the weights in a neural network just as long as they’re randomly initiated!

In particular, we train both a bigram neural network language model and a trigram neural network language model. In both cases, we use word vectors of length 10 meaning that the input layer of the bigram NN is a 10 (a single word) and the input layer of the trigram is 20 (2 words concatenated together). This input layer then maps to the single hidden layer with 10 hidden units through a tanh activation function. Finally, the hidden layer is then connected to the output layer which is has size $|V|$ through a softmax activation function. We use stochastic adaptive gradient descent so that the errors can be properly backpropagated to the word vectors. We purposefully keep our network relatively small since our data set is rather large and training even this small neural network takes considerable time.

Unfortunately, our neural networks are still in the process of training. Due to various complications, we were not able to train our model for as long as we would’ve liked. Since our training corpus contains over 3,000,000 observations, each complete pass through it takes approximately 4.5 hours. At the time of writing this section, we’ve only completed 4 full passes.

We adapted our code so it saved a useable snap-

shot of the network after each full pass through the data. Unfortunately, as a result of an oversight in the code, we forgot to store the NALL of training data in this snapshot. All is not lost since we remembered to at least store the NALL on the validation corpus. We are highly confident our neural network is working as intended since we see that the NALL decreases every pass on the validation corpus in Figure 10 below.




Figure 9. The negative average log-likelihood of the bigram and trigram models

We fully expect the validation corpus NALL to eventually start increasing once the neural network begins to overfit the training corpus, however we simply haven't reached that point yet.

Despite the fact that our training isn't finished, both our neural network-based language models have already achieved lower negative average log-likelihood compared to the smoothed standard bigram model. We build 2 sets of article perplexity scores. One set is derived from the smoothed bigram model to serve as a base comparison. Our other set is derived from the most recently completed pass of our trigram NN language model. We report the following article perplexity scores distribution:

It's important to note however, that we only care about the goodness of the underlying language model insofar that a better language will produce features with better predictive power. However, it could very well be the case features generated by a worse language model can work just as well as features generated by a better language model in




Figure 10. The Distribution of Bigram and trigram Neural Network Perplexities

the regression model if the relative variation of articles is maintained. Hence our reasoning to produce the two sets of perplexity scores.

4. Predictive Regression Model

Using our full set of features, we are now able to perform the regression task we originally had in mind, and attempt to determine how (if at all) content drives viewership. We regress $\log(\text{article pageviews})$, y , on our design matrix, Φ , which includes entries for each of our $k - 1$ features, plus an intercept term (in this case, $k = 102$). We estimate the feature weights using the closed form solution for OLS and ridge regression:

$$\beta = (\Phi^T \Phi + \lambda \mathbf{I})^{-1} \Phi^T y, \quad (7)$$

where \mathbf{I} is the $k \times k$ identity matrix, and λ is our regularization parameter. Setting $\lambda = 0$ corresponds to OLS, whereas a non-zero value of λ corresponds to ridge regression. The motivation for performing ridge regression as opposed to OLS is to not overfit on our data, and the value of λ can be interpreted as the strength of our Bayesian prior on the feature weights being equal to 0 [6].

In order to choose an appropriate value of λ , we split our data into training, validation, and test sets. 90% of the data is allocated to the training set, 10% to the validation set, and 10% to the test set. Although we estimate β on the training data using the above closed-form solution, we cross-validate on our validation set for each λ , and choose

the value of λ that produces the lowest MSE on our validation dataset. To ensure that we have not overfit β to our validation dataset, we also calculate the MSE on the test set as a final step. A comparison of the training, validation, and test MSEs for various values of λ is found in Figure 13.

We find that $\lambda = 100$ minimizes the MSE on our validation set. Table 6 displays the 20 feature weights with the largest magnitudes $\lambda = 100$. Two charts showing weights for the full set of features (excluding the intercept term) can be found in Figures 11 and 12.

Table 6. 20 Most Significant Weights

Feature	Weight
intercept	8.964
log(Word Count)	0.741
Desk: Foreign	0.550
Desk: Travel	-0.465
Section: World	-0.402
Section: Opinion	0.283
Type: BlogPost	-0.282
Type of Material: Schedule	-0.267
Desk: None	-0.262
Section: Movies	-0.239
Time of Day: 12-17	0.232
Type of Material: Review	-0.218
Desk: National	0.211
Section: Health	0.210
Section: Books	0.201
Type of Material: Letter	-0.200
Type of Material: News	0.187
Section: Sports	-0.175
Type of Material: Op-Ed	0.170
Desk: BookReview	-0.154

It's worth taking the time to discuss Figure 13, which shows the training and holdout MSE for various values of λ , in slightly more depth. There are a few things in this plot worth discussing. First, note that the validation MSE is consistently higher than the training data MSE, which is consistently higher than the test data MSE. Given the (relatively) small size of our dataset (6,687 observations), this is likely due to the sampling we used to separate our

data into training, validation, and test data. However, we don't expect this effect the validity of our cross validation.

Another thing worth noting is that even on the training dataset, there exist non-zero values of λ that achieve a lower MSE than the OLS estimate of β . At first, this was surprising to our group, as conceptually OLS is often thought of as the linear regression method that minimizes MSE. However, it is important to note that OLS only holds this distinction amongst unbiased estimators. Hoerl and Kennard (1970) [7] prove the existence theorem for ridge regression, which claims the existence of some λ such that β_{ridge} produces a lower MSE than β_{OLS} .

Another way of framing this finding is through bias-variance tradeoff. Recall that the MSE can be written as a function of the bias and variance:

$$MSE = (\text{Bias})^2 + \text{Var.} \quad (8)$$

For some values of λ , ridge regression is able to lower the MSE by decreasing variance, but increase the bias from zero to some non-zero value. We believe the changes in MSE we observe in our dataset as we vary λ can be explained by this phenomenon.




Figure 13. The effect of regularization on MSE

In general, we can now interpret the strength of the feature weights produced by our regression to determine how predictive a particular feature is of readership. Note from Table 6 that most of the features with the most predictive power are not the text-based features. The intercept term in our regression is orders of magnitude larger than any other




Figure 11. Linear Regression Feature Weights (excluding intercept term) (1 of 2)

feature, implying that most of the NYTimes articles receive many pageviews in the baseline case. The strongest coefficients tend to be those that indicate the publication desk, section, and time of publication of the article. This suggests that the content of an article itself may not be as important as the context in which it is published. The one exception we see is that a higher word count is predictive of higher viewership. We suspect that what's going on here is correlative, rather than causal - the quality of longer pieces (e.g., NYTimes Magazine articles) is likely higher, thus driving more readership. But

we don't expect that a website full of low-quality, 5,000 word pieces would be successful.

4.1 The Effect of Textual Features

Given the large amount of work we put into building numerous text-based features (such as the trigram neural network perplexity, the sentiment labels, and the Flesch reading ease) and the relatively low impact they seem to have had on our regression (based on feature weights), we want to specifically evaluate the impact of these features on our regression. Specifically, how much incremental reduction in




Figure 12. Linear Regression Feature Weights (excluding intercept term) (2 of 2)

MSE are we getting by including them? We first re-run the exact same regression specification, but instead of using the trigram perplexity calculated from a neural network, we instead use the simple bigram perplexity discussed earlier in this paper. We find that this actually leads to a reduction in MSE, from 2.392 to 2.389. This small, but modest change suggests that there is currently almost no incremental value from using a neural network language model as opposed to a more straightforward language model.

As a next step, we ask if text features in general

add much value to our model. Given that, in general, the magnitude of text feature weights is dwarfed by the magnitude of contextual feature weights, we might expect that text features do not add much. We find that removing the perplexity, sentiment, reading ease, and word count features from our regression leads to an increase in MSE, from 2.392 to 2.502. This result is encouraging, as it suggests that even if text features currently aren't contributing as much as we had hoped, they are doing something! Removing them from our model leads to a 4.6% increase in the MSE.

Table 7. Comparison of MSEs for different model specifications

	NN Trigram	Bigram	No Text
Best Val. MSE	2.392	2.389	2.502
Train. MSE	2.011	2.011	2.274
OLS Val. MSE	2.407	2.405	2.506

5. Future Work

Unfortunately, we were unable to implement all of several important features we initially wanted to include in our regression, due to a combination of time constraints and various unanticipated technical difficulties and errors. Most notably, this paper currently lacks some form of topic modeling, which we expect to be a strong predictor of content viewership (at least in our context). We also suspect that the addition of features containing information about article headlines may provide some improvement. It seems reasonable to assume that an article's headline is a fairly important factor in a user's decision to click through or not.

There is also room for improvement in the design and extraction of our existing features. For example, the perplexity score feature ultimately provided very little predictive power. This may be due to the issues in the quality of the underlying language model. This is particularly likely given that the neural network model was not able to complete many iterations, and has not yet converged. Although we observed virtually no difference in performance moving from the bigram perplexity scores to the trigram neural network perplexity scores, this again may be due to insufficient training time for the neural network. We find it conceivable that, given many more cycles to train, our neural network would start to yield visible performance gains over the bigram perplexity score. Another hypothesis is that rather than measuring the the perplexity of an article relative to the entire corpus of articles, it would be more effective to measure a given article's perplexity relative to other articles that cover the same event or topic.

There is ample room to improve our sentiment analysis methodology. The training dataset gener-

ated using mechanical turk is relatively small (200 training articles). Because of this, there may exist many words that have a strong probability of appearing in an article conditional on sentiment that simply do not appear in our training set. With more time (and money!), we could label more data and improve the accuracy of our model. Furthermore, the model currently skips words that do not occur in the training corpus. An extension of our Naive Bayes implementation could use a method such as Laplace smoothing, so as to not simply ignore words we haven't seen in our training data.

In addition, the current discrepancies between our implementation of Naive Bayes and NLTK implementation of the algorithm, while not large in magnitude, are alarming. In the future, we hope to dig deeper into this discrepancy and identify the root cause. Our current hypothesis is that NLTK's implementation of Naive Bayes does not take into account the number of times a particular word appears in a document. Because of this, a larger number of distinct informative words are required in order to overcome the strong prior belief that a given article's sentiment is neutral.

Another avenue for potential improvement to our model is the application of basis expansion to our variables. This would allow us to include polynomial and interaction terms. In many cases, there is no good justification for assuming a feature is linearly related to the output feature. Hence, applying a polynomial expansion might uncover an entirely different relationship between our dependent variable and its covariates. By a similar line of reasoning, we suspect that interaction terms may have significant predictive power.

Given that we have access to six full months of NYTimes data, we also hope to expand our dataset to include the full body of NYTimes articles we have available to us. With a larger sample size and more expansive corpus, we might expect to see drastically different optimal weights, as our results will be less sensitive to time-dependent trends in viewership and content. Since we also have access to Twitter data tracking every tweet and retweet involving a NYT article, we also hope incorporate this data into a future version of our model. This

would enable us to explore potential cross effects, such as the relationship between article content and Twitter sharing.

6. Division of Labor

The work for this project was divided as follows. Michael extracted the New York Times contextual data and page view counts and wrote the NYT website scraping script. Dave and Jeremy built out many of the content-driven features and additional contextual features, such as word count, author gender, popularity, and Flesch reading ease. Dave and Jeremy built the implementation and validation of the Naive Bayes classifier, while Michael focused his efforts on the construction and validation of the language model. Dave, with assistance from Michael, focused on the implementation of OLS and ridge regression, along with validation and discussion. Plot generation, table creation, writing, and editing was evenly split between Michael and Dave. Extenuating personal circumstances unfortunately prevented Jeremy from contributing to the project as much as originally anticipated - he will likely be in touch soon with the instructors regarding this issue.

Acknowledgments

We would like to thank Professor Sinan Aral for his work procuring this dataset, as well as his limitless advice, mentorship, and patience. We'd also like to thank Paramveer Dhillon for his useful comments and his work pre-pre-processing our dataset.

References

- [1] Jonah Berger and Katherine L Milkman. What makes online content viral? *Journal of marketing research*, 49(2):192–205, 2012.
- [2] Rudolph Flesch. A new readability yardstick. *Journal of applied psychology*, 32(3):221, 1948.
- [3] Jason D Rennie, Lawrence Shih, Jaime Teevan, David R Karger, et al. Tackling the poor assumptions of naive bayes text classifiers. In

ICML, volume 3, pages 616–623. Washington DC), 2003.

- [4] Greg Lamp. Naive bayes in python. <http://blog.yhathq.com/posts/naive-bayes-in-python.html>. Accessed: 2015-12-06.
- [5] Regina Barzilay and Tommi Jaakola. 6.864 lectures 2 and 3 notes. 2015.
- [6] Guy Bresler and Tatsunori Hashimoto. Lecture notes on: regularization and bias-variance trade-offs, September 2015.
- [7] Arthur E Hoerl and Robert W Kennard. Ridge regression: Biased estimation for nonorthogonal problems. *Technometrics*, 12(1):55–67, 1970.